Injestion Engine README

Scope

Hard-coded values should be moved into the object configuration so the engine can be reused without editing internal logic.[web:1062]
The interface should be created after the object implementation is complete, which keeps the contract aligned with the final behavior of the engine.[web:1070][web:1080]

Project flow

The injestion engine is being built in stages so document intake, AI enrichment, storage, and querying are separated into clear responsibilities.[web:1062][web:1078]

Stage 2

The engine will take all documents in the `__injest` directory and process them as the working input set.
It will break files down into their readable content such as images, text, tables, headers, and footers.
It will then feed one large string object into a local LLM for processing, which is appropriate for sensitive work that should not be sent to a cloud service.

Stage 3

The AI output will be captured and embedded into a vector database using sample files as the initial reference set.
A sample reference video for this stage is: YouTube sample.

Stage 4

The front end will allow a user to submit a query.
The backend will receive that query, search the database, and return the matching object or objects.

Notes

- Keep configuration values out of method bodies where possible.
- Finish the main object first, then create the interface after the final object shape is stable.
- Keep the implementation ordered in the same top-down sequence as the controller flow for readability.[web:1070][web:1077]