

# Vector Databases

Group 3

Cristi, Favour, Vini, Yung



# Outline

**1: What is a vector database**

**2: Why is the technology relevant?**

**3: How does it work**

**4: What are some alternatives?**

**5: What is a use cases for a Vector database?**

**6: Demo of code**

**7: Quiz: Introduction to Vector Databases**

# 1: What is a vector database?

**Vector:** quantity with both **magnitude(size & length)** and direction

Words have magnitude and direction:

“it’s not bad”

->

<-

“it’s not good”

“i like it”

—>

←-

“I dislike it”

“i love it”

----->

←-----

“I hate it”

“I’m obsessed”



----->

←-----

“i despise it”

Example of high dimensional data:

**“I dislike dogs” VS “i love cats” VS “fish is tasty”**



Source: Douglas  
College, 2025

## 2: Why is the technology relevant?

**Data is enriched by being compared in higher dimension.**

**1.Semantic search and similarity matching:** find content based on meaning and not key words.


- **Example:** “House price” <- “I like animals ” -----> “i like cats”

**2.Nearest neighbour search for recommendation systems :** find similar data points to cluster.

- **Example:** “i like lons” <--- “i like felions“ -----> “I like cats”

**3.Retrieval Augmented Generation (RAG):** Make the LLM provide answers based on relevant embedded/vectorized documents.

- **Example:** “find a document on felions” -> documents on cats/lions/panthers/....




**Source:** Wolfram 2014

# 3: How does it work?

**Chunk overlap:** semantics are captured across clusters of words.

**Embeddings:** Semantic encoders extract meaning from words.

**Scoring:** Calculate the relevance score of the query vs existing documents.



Source: Cristi S, 2026


# 3A: Chunk overlap

100 page document is chunked into a page.

Each page is further chunked into paragraphs, then sentences.


Each sentences is chunked into parts based on meaning and intensity.

Chunks overlap to retain inter-chunk context.



Source: Cristi S, 2026

# 3B: Embeddings



A diagram showing an 'Embedding engine' box with an input arrow from above and an output arrow pointing to a table of numerical data.

|        |          |          |          |          |          |          |          |          |          |          |     |
|--------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|-----|
| PDF    | 0.044963 | 0.034968 | 0.866643 | 0.330055 | 0.436304 | 0.143085 | 0.932044 | 0.051222 | 0.631056 | 0.264502 | ... |
| Word   | 0.827491 | 0.810262 | 0.229789 | 0.493483 | 0.731988 | 0.514412 | 0.422001 | 0.173519 | 0.173632 | 0.753861 | ... |
| Excell | 0.798445 | 0.728105 | 0.890631 | 0.636032 | 0.819301 | 0.136202 | 0.946666 | 0.894651 | 0.720499 | 0.863904 | ... |
| Access | 0.77067  | 0.940332 | 0.368095 | 0.899787 | 0.044978 | 0.055753 | 0.853838 | 0.259878 | 0.656708 | 0.69291  | ... |


Source: Cristi, 2026

**Transformers process chunks into vectors using self-attention mechanism for context.**

- **Example:** “Bank account” vs “river bank”

**Neural Networks like BERT can be used to create embeddings.**

**Vectors are scatters on the 3d coordinate grid**



Source: Desmos, 2026



# 3B: Scoring

**Storage and Retrieval is based on finding nearest neighbour vectors that share vector direction.**

**A query is embedded and its vectors are scattered, finding the nearest neighbours.**

**Vectors can capture momentum of a topic based on length of vector.**

- **Example:** 15, 25, 150 users asking for cat images.



20° High similarity  
90° No similarity  
180° Perfectly opposite

**Source:** Cristi, 2026



# 4: What are some alternatives?

**ChromaDB:** Open source database that uses LangChain and LlamaIndex to support virtualization.

- **Key feature:** Ease of use.

**Pinecone:** close source that excels at high-dimensional data. Excels at machine learning applications.

- **Key feature :** performance at scale.

**Weaviate:** Open source AI native vector database, which allows the storage of data objects and vector embeddings. Excels at searching vectors.

- **Key feature:** store objects with metadata, properties, and relationships for complex data rich search.



**Source:** Chroma , 2026



**Source:** Pinecone, 2026



**Source:** Weaviate, 2026

# 5: Use case for vector databases.

## Problem

Traditional databases and keyword-based search systems retrieve information mainly through exact term matching. This approach may fail when a query and a relevant document express the same concept using different vocabulary.

For example:

Query: “**How can I reduce cloud infrastructure costs?**”

Document: “**Strategies for optimizing computing resource expenditure**”

Although semantically related, the texts share few exact keywords.

Traditional keyword search may fail to retrieve all relevant listings because it depends on lexical overlap rather than semantic similarity.

By converting each property listing into a vector **embedding**, the system performs **semantic search**, retrieving results based on meaning rather than exact keywords.

# 5: Use case for vector databases.

## Problem

Organizations generate thousands of unstructured marketing assets, including:

- Campaign briefs
- Product descriptions
- Email campaigns
- Presentations
- Brand guidelines
- Social media content

Traditional keyword search may fail when relevant materials express the same idea using different terminology.

By converting each file or text into a vector **embedding**, the system performs **semantic search**, retrieving results based on meaning rather than exact keywords.

# 5: Use case vector databases.


example

If there's a query "Sustainability campaign" .

Documents may instead contain:

- Green initiatives
- Eco-friendly products
- Environmental responsibility
- Carbon reduction strategy

Although semantically related, exact keyword search may miss relevant documents.



# 5: Use case for vector databases.

## Marketing Content Pipeline

Embedding Representation

$x_i \rightarrow e_i \in \mathbb{R}^d$

where

- $x_i$  = document chunk
- $e_i$  = embedding vector
- $d$  = embedding dimension

Each document is transformed into a numerical representation that captures semantic meaning.

# 5: Use case for vector databases.

## Semantic Retrieval

The user query is encoded using the same embedding model:

$$q \rightarrow e_q \in \mathbb{R}^d$$

The vector database compares the query embedding against all indexed vectors using similarity metrics such as:

- Cosine Similarity
- Euclidean Distance
- Dot Product

Approximate Nearest Neighbour (ANN) indexing enables efficient similarity search over millions of vectors.

## Retrieval Function

$$R(q) = \text{TopK}_{e_i \in D} [\text{sim}(e_q, e_i)]$$

The system returns the **Top-K** most semantically relevant documents.

# 5: Use case for vector databases.

## Business Benefits



- Faster document discovery
- Reduced duplicated work
- Knowledge reuse
- Context-aware search
- Better decision support
- Improved productivity

## Example Query

"Find previous sustainability campaigns targeting young consumers."

### Returned Results

- ✓ Eco-Friendly Summer Campaign
- ✓ Green Product Launch
- ✓ ESG Brand Strategy



# 5: Use case for vector databases.

## Real-World Applications

- Marketing asset management
- Enterprise knowledge bases
- Customer support systems
- Retrieval-Augmented Generation (RAG)
- Document recommendation systems
- Intelligent enterprise search



# 6: Demo code-ChromaDB Vector Search



**Semantic search — find by meaning, not by keyword**

**What we're demonstrating:**

**1. Ingest Documents — Upload course PDFs / PPTX into the UI**

**2. Embed & Store — Convert text into vector embeddings in ChromaDB**


**3. Semantic Search — Ask in natural language, retrieve by meaning** Key idea: **Keyword search finds words. Vector search finds meaning.**

# 6: Step 1: Ingest Documents

1. Upload course PDFs / PPTX into the UI

2. Extract text from the files (OCR for scanned documents)

3. Documents enter the pipeline, ready for embedding





# 6: Step 2 — Embed & Store

**1.Chunking: split long documents into meaningful segments**

**2.Each chunk is converted into a vector embedding (via Ollama)**

**3.Vectors are stored in ChromaDB — similar meanings sit closer together**



# 6: Step 3 — Semantic Search

1. Ask in natural language — no exact keywords needed

2. The query becomes a vector; ChromaDB finds the closest matches

3. Distance score = semantic relevance (smaller = more relevant)

**VanCity Rental**  
Market intelligence workspace

**Reports**  
Review summaries, analytics output, and exported insights.

Updates **1** Logout

- Dashboard
- Gmail
- Documents
- Spiders
- Reports**
- Feedback
- Settings

Use Ollama to summarize uploaded housing documents and review generated report output in the Reports page.

what is a vector database

Open AI Generated Report

**2**

**AI Answer**

A vector database is a type of database that stores and manages data in a form called vectors, which are mathematical representations of objects. These vectors capture the magnitude (size or length), direction, and words associated with each piece of data. \*\*Source\*\* Document 1: Title "realestate rental app big data 26" Key points: - Vector databases store data as vectors for efficient similarity searches and semantic analysis. - It enables retrieval based on meaning instead of keyword search alone.

**Matched documents**

**realestate rental app big data 26**

Overview of the document content about realestate rental app big data 26

Source: upload | Type: Unknown

Sender: Unknown | Date: Unknown

**History**

what is a vector database  
2025-07-22 21:51:37

**"distances":** [

- 0.41436338424682617,
- 0.780095100402832,
- 0.8720766305923462,
- 0.9028581976890564,
- 0.9243981838226318

]

]

best match

# 7: Quiz: Introduction to Vector Databases

Ready? Scan now and let's start!



## 2. Choose Your Format

Groups can choose one of the following three tracks for their seminar:


- **The Research Deep-Dive:** Analyze a recent peer-reviewed paper (published within the last 24 months) regarding Big Data architectures, AI integration, or data ethics.
- **The Technology Spotlight:** Introduce a new or emerging tool/framework (e.g., DuckDB, Polars, dbt, or Vector Databases) and explain its position in the ecosystem.
- **The "Flash Course" (Hands-on):** A practical, tutorial-style session teaching the class how to use a specific library or tool to solve a data problem.

## 3. Mandatory Content Requirements



- **The "Why" Factor:** Why is this research/technology relevant right now?
- **Technical Architecture:** A clear explanation of how the system/concept works internally.
- **Comparative Analysis:** How does this compare to existing solutions (e.g., Spark vs. Flink)?
- **Interactive Element:** You must include a Q&A, a live poll, a short quiz, or a code demo to engage the audience.

## 4. Deliverables

- **Presentation Slides** (PDF or PowerPoint).
- **Seminar Summary** (1-page handout for classmates).
- **Resource List** (links to documentation, papers, or Github repos).

A decorative grid pattern of squares, some solid light gray and some outlined in white, arranged in a stepped fashion on the left side of the slide.

**Thank you**  
**Questions?**

A small decorative grid pattern of squares, some solid light gray and some outlined in white, located in the lower right area of the slide.A large, solid yellow curved shape, resembling a stylized hill or a large arc, positioned in the bottom right corner of the slide.